

China, Pakistan, Myanmar and the Belt and Road

2026-08-26 • 34 Sources

34

REGISTERED SOURCES

34/34

SUPPORTED CLAIMS

20,985

AIDDATA GLOBAL
PROJECTS

5

MAJOR FINDINGS

EXECUTIVE SUMMARY

This report defines a defensible way to make Palimpsest and NarcoScope a deep public-evidence system for Chinese economic activity in Pakistan and Myanmar, including CPEC, Gwadar, the China-Myanmar corridor, Kyaukphyu, and the aggregate historical context of conflict in Balochistan. Its central conclusion is that scale should come from traceability, not accumulation alone. AidData's global dataset lists 20,985 Chinese-financed projects across 165 low- and middle-income countries, while its geospatial companion covers 9,405 projects across 148 countries [4] [34]. Those totals show why a durable source registry, revision history, and project-identity model are indispensable.

The economic case must remain conditional. The World Bank model estimated that completed Belt and Road transport corridors could reduce travel times by up to 12 percent, but that is a modeled corridor effect rather than a measured return for Gwadar, Kyaukphyu, Pakistan, or Myanmar [3]. Pakistan's own sources present Gwadar, energy, transport, and industrial cooperation as priorities; official pages also list completed and ongoing assets [1] [2] [16] [17] [22]. At the same time, Pakistan's public external debt and liabilities reached US\$103.7 billion at end-FY25, and the IMF describes repayment capacity as dependent on policy implementation and timely external financing [20] [21]. The system therefore must connect each infrastructure assertion to financing, construction status, trade, macroeconomic exposure, and source confidence without implying causation.

Myanmar requires an even stronger status model. Authorities continue to describe Kyaukphyu cooperation and governing agreements, while the Asian Development Bank reports that its public-sector disbursements and new contracts have been on hold since February 2021 [23] [24] [26]. “Announced,” “contracted,” “under review,” “paused,” “built,” and “operating” cannot be collapsed into one project state.

Balochistan coverage should explain historical periods, actor continuity, rights concerns, and the relationship between insecurity and infrastructure at an aggregate level [11] [28] [29] [30] [31] [32]. It must exclude tactical routes, live locations, personal profiles, and targeting-relevant detail. The recommended architecture is a rights-gated Palimpsest evidence ledger feeding a deliberately coarser NarcoScope publication layer. Unknown, stale, disputed, or unlicensed material remains visibly unavailable rather than becoming an apparently neutral zero.

INTRODUCTION

Research question and purpose

The research question is: how can Palimpsest and NarcoScope become a reliable, deeply integrated public-evidence system for China's work in Pakistan and Myanmar, the Belt and Road's infrastructure and economic dimensions, CPEC and Gwadar, the China-Myanmar corridor and Kyaukphyu, and the aggregate historical and political context of the Balochistan movement, while preserving publication rights and avoiding operational intelligence?

The intended audience is broad but demanding: researchers, journalists, policy analysts, civil-society organizations, investors, and public officials who need to understand not only what a project page says, but who said it, when it was retrieved, what version was observed, whether it can be republished, what independent evidence exists, and what remains unknown. The report is therefore both an evidence synthesis and an implementation specification. It does not claim that the 34 registered sources constitute an exhaustive inventory of every Chinese project, financing instrument, or conflict event. It defines the deepest product that can be responsibly supported by the evidence currently stored.

Scope

The infrastructure scope includes official CPEC and Gwadar portfolios, Pakistan's macroeconomic and planning records, the joint CPEC long-term plan, Myanmar's official descriptions of Kyaukphyu and special economic zones, and independent or multilateral sources on corridor economics and development finance [1] [2] [3] [4] [15] [16] [17] [18] [19] [20] [21] [22] [23] [24] [25] [26] [27] [34]. The data-system scope includes source identity, project identity, version history, evidence spans, rights state, freshness, uncertainty, publication transforms, and query behavior. UN Comtrade is included as a future trade-measurement input, with the important constraint that its preview endpoints are keyless but limited in records and rate [33].

Conflict coverage is intentionally bounded. UCDP provides freely available CC BY 4.0 datasets, versioned replication packages, conflict-year records, and stable actor identifiers [7] [28] [29] [30]. ACLED documents both coverage enhancements and known limitations, but its EULA governs what users may do with its data and content [9] [10]. These sources can support aggregate time series, documented actor continuity, uncertainty ranges, and historical context. They do not justify publishing event-level coordinates, descriptions that expose individuals, or rapidly updating operational feeds. The Balochistan section consequently focuses on historical sequence, institutional grievance, rights context, and the analytical relationship between infrastructure and insecurity.

Method and source treatment

Each source was registered under a stable source identifier, and each persisted evidence row records a locator, evidence type, retrieval query, capture time, and short quote or paraphrase. Display citations [1] through [34] follow registration order; the underlying source identifiers do not change when prose is

reordered. Official Chinese, Pakistani, and Myanmar materials are treated as attributed primary claims. Multilateral datasets and academic work are treated as independent evidence with their own model, sampling, and licensing constraints. Human-rights reporting is used for contextual risk and attribution, not as a substitute for a conflict event database [31].

Four assumptions govern the analysis. First, public evidence is necessarily incomplete, so absence is not evidence of no activity. Second, infrastructure benefit depends on complementary policy and operating conditions; an announced or completed asset is not itself proof of economic additionality. Third, official project classifications can be useful without being neutral. Fourth, safety and rights are data-model requirements rather than editorial footnotes: records that are unlicensed, stale, disputed, or too granular for safe publication remain quarantined.

MAIN ANALYSIS

Finding 1: Belt and Road economics is a chain of conditional effects, not a single project total

The strongest economic framing in the current evidence is the World Bank's corridor model. It estimated that travel times could decline by up to 12 percent once the modeled Belt and Road transport corridors were completed [3]. That statement is useful, but it has three boundaries. It is a forecast produced by a model; it applies to a network of corridors rather than one port; and it says nothing by itself about how benefits are distributed among countries, regions, firms, or households. Palimpsest should preserve all three boundaries as fields attached to the claim: `measure_type = model_estimate`, `geographic_scope = corridor_network`, and `project_attribution = not_established`.

This matters because infrastructure narratives often jump from physical completion to economic success. A port can be built without attracting the forecast cargo; a road can lower engineering travel time without lowering border delay; an energy plant can add nameplate capacity without supplying reliable power at an affordable tariff. None of those failure modes proves that a project is harmful. They show that “built” and “beneficial” are separate propositions requiring separate evidence. The BRI Trade Costs Database is explicitly public and CC BY 4.0, making it suitable for reproducible baseline and scenario inputs [14]. Future benefit claims can therefore show their assumptions and calculation lineage instead of appearing as free-standing promotional figures.

Scale also creates an identity problem. AidData's Global Chinese Development Finance Dataset lists 20,985 projects across 165 low- and middle-income countries [4]. The TUFF methodology describes version 3.0 as covering 165 countries, with identified projects in 146 [13]. Its geospatial product reports coverage of 9,405 projects across 148 low- and middle-income countries [34]. These figures should not be merged into a single “total BRI projects” number. They refer to related but distinct products, inclusion rules, and geographic completeness. The correct system behavior is to retain the dataset version, observation unit, coverage definition, and publication date alongside every count.

The resulting economic evidence chain should contain at least five linked but independently testable layers. The first is the asset layer: port, airport, road, rail, power, special economic zone, digital link, or municipal service. The second is the finance layer: commitment, loan, grant, equity, guarantee, sovereign exposure, currency, maturity, and disbursement state. The third is the delivery layer: announcement, agreement, financial close, construction, completion, operation, suspension, cancellation, or unknown. The fourth is the usage layer: throughput, capacity utilization, travel time, outages, occupancy, firm registrations, or trade flows. The fifth is the outcome layer: productivity, prices, fiscal burden, employment, distribution, environmental effects, and resilience. The sources presently support only fragments of this chain, and the gaps must remain visible.

Trade measurement is a particularly important future layer because it can test whether corridor operation is associated with observable changes. UN Comtrade's preview APIs require no key but impose record and rate limits [33]. That makes them useful for small reproducibility checks, not an unlimited production collector. A mature collector should record the reporter, partner, commodity classification, flow, period, revision date, API tier, and extraction query. It should also distinguish nominal trade value from volume and avoid claiming that contemporaneous growth was caused by BRI infrastructure. A credible dashboard can show “trade changed after opening” while withholding “the project caused the change” until an identification strategy supports it.

The overall economic conclusion is therefore conditional rather than negative or promotional. Corridor infrastructure can reduce a real friction, and the World Bank model gives a bounded estimate of one such mechanism [3]. AidData demonstrates that Chinese development finance can be studied at a global project scale [4] [13] [34]. Publicly licensed World Bank inputs enable transparent replication [14]. Yet no current evidence row establishes the realized return of Gwadar or Kyaukphyu. The product should reward that honesty: a well-populated project page with an unresolved outcome is more reliable than a confident score built from incomparable inputs.

Finding 2: CPEC and Gwadar need a project ledger connected to Pakistan's macroeconomic constraints

The joint CPEC long-term plan identifies Gwadar port, energy, transport infrastructure, and industrial cooperation as four priorities [22]. Pakistan's CPEC portal separately lists Gwadar projects and a broader progress inventory, including a 1,320-megawatt China Hub coal power project in Balochistan [1] [2]. These official sources are indispensable for discovering project names and government classifications. They should appear in the system as issuer-attributed claims with a retrieval timestamp and page revision, not as self-validating proof of current operation.

Gwadar's official records already illustrate why project identity is necessary. The Gwadar Port Authority lists 17 completed and three ongoing development projects [17]. Its free-zone page says that more than 20 enterprises are registered [18]. Its master-plan page states that concession rights transferred to China Overseas Port Holding Company in 2013 [19]. The Pakistan Annual Plan 2025–26 lists the New Gwadar International Airport as completed [16]. Those are four different observation types:

portfolio count, enterprise registration count, concession event, and asset status. They can be related through the Gwadar geography and CPEC program, but they must not be collapsed into a single progress percentage.

A trustworthy Gwadar page should therefore show a project family rather than one blended card. Each asset would have a canonical project identifier, known aliases, sponsor and implementing-authority claims, a status timeline, finance instruments, physical components, and evidence spans. The system should also record contradictions. If an authority calls an asset completed while an independent source reports limited operation, both observations remain, each with its date and scope. A derived status may be published only if its rule is transparent—for example, `completed_physical` can coexist with `operational_evidence_unavailable`.

Macroeconomic context must sit beside, not inside, project status. The Pakistan Economic Survey 2024–25 publishes distinct chapters for public debt, trade, transport, energy, and investment [15]. The IMF country page identifies an Article IV country report dated October 10, 2024 [6]. The State Bank of Pakistan reported US\$103.7 billion in public external debt and liabilities at end-FY25 [20]. The IMF's second-review material says repayment capacity remains dependent on policy implementation and timely external financing [21]. These facts justify monitoring debt composition, external financing needs, and balance-of-payments conditions. They do not establish that CPEC caused Pakistan's entire debt stock, nor do they identify the share attributable to any one Chinese project.

This separation enables better questions. A user could ask which Gwadar projects an official authority currently calls completed; which have a concession or governing agreement; which have reported operating or usage evidence; which have finance data; and which sit in a period of elevated external-financing risk. A user should not receive a synthetic “debt trap” or “development success” label unless the underlying measures and causal reasoning are made explicit. The product's role is to expose the evidence chain, not encode a geopolitical slogan.

The 2017 USIP report described the Baloch insurgency as the greatest roadblock to developing Gwadar port [11]. That is a dated analytical judgment, not a timeless quantitative result. It remains relevant because it shows that security, political legitimacy, and infrastructure delivery are interdependent in published analysis. It should be stored with `publication_year = 2017`, `claim_type = analysis`, and an explicit stale-context warning when used to interpret later developments. Current conflict or rights conditions require newer sources and should never be inferred from a nine-year-old assessment alone.

The deepest useful Gwadar product would combine three views. A provenance view would display every source, evidence span, observed date, rights status, and revision. A project view would show assets, agreements, financing, delivery state, and operation. A context view would present macroeconomic conditions, aggregate conflict history, human-rights reporting, and data gaps. Keeping these views linked but distinct prevents the two most common errors: attributing national macro conditions to a single corridor without evidence, and treating physical completion as proof of broad local benefit.

Finding 3: Myanmar corridor coverage must represent legal continuity, implementation interruption, and economic deterioration separately

The current Myanmar evidence contains a real tension. The Myanmar Port Authority describes cooperation with the CITIC consortium for the Kyaukphyu deep-sea port [23]. A January 2025 Myanmar Ministry of Information account says that a 2020 concession agreement and a 2023 addendum governed the project's timetable [24]. The Directorate of Investment and Company Administration identifies Thilawa, Kyaukphyu, and Dawei as Myanmar's special economic zones [25]. These official statements establish that authorities describe a legal and institutional project structure. They do not, on their own, establish construction progress, financing disbursement, operational throughput, or local economic outcomes.

The Asian Development Bank supplies a different observation: public-sector project disbursements and new contracts have been on hold since February 2021 [26]. This does not prove that every Chinese or privately financed project was frozen, because the statement concerns ADB's own public-sector activity and contracting. It does show that a major development institution treats the post-2021 operating environment as materially discontinuous. A project database that labels Kyaukphyu simply “ongoing” would conceal this institutional context; one that labels it “cancelled” would claim more than the evidence supports.

The necessary lifecycle is multi-dimensional. `legal_state` records concession agreements, addenda, approvals, and implementing entities. `finance_state` records commitments, disbursements, equity, guarantees, and unknowns. `physical_state` records construction evidence and completion. `operating_state` records throughput, service delivery, or non-operation. `context_state` records macroeconomic-monitor availability, conflict aggregation, and institutional interruptions. Each dimension has its own observed date and source. A public summary can then say, for example, that official sources describe a governing agreement while independently verified physical or operating evidence remains unavailable.

The World Bank's Myanmar Economic Monitor series offers a biannual view of economic developments and key challenges [5]. The Asian Development Outlook April 2025 is released under a Creative Commons Attribution 3.0 IGO license [27]. These sources can anchor macroeconomic series and narrative context, but their observations must remain edition-specific. An index should point to the exact report, table, unit, reference period, release date, and revision rather than copying a number into a permanent “Myanmar economy” field. In a volatile economy, observation time and publication time are different clocks.

The same discipline applies to security data. ACLED says its Myanmar coverage spans January 2010 to the present [12], but “to the present” is not a guarantee of complete real-time observation. Its broader methodology documents enhancements and known limitations, and reuse is controlled by its EULA [9] [10]. UCDP's licensed versioned datasets may provide a clearer publication path for aggregate historical indicators [7] [28]. The product should expose each dataset's period, update cadence, methodology, and rights state rather than fusing event counts from different systems.

Myanmar therefore should not be a copied version of the Pakistan module. Both use the same evidence primitives—source, claim, project, finance, status, geography, and time—but Myanmar needs stronger discontinuity semantics. It also needs stricter separation between authority statements and independent observation because sources operate within a contested political environment. The system can remain useful without choosing one political narrative as ground truth: it can show who asserted what, when, under which legal instrument, and what corroboration is absent.

Finding 4: Balochistan coverage should be historically comprehensive but operationally coarse

“The Balochistan liberation movement” should not be modeled as one timeless actor, one uninterrupted campaign, or one explanation for every event. The academic study registered here examines the evolution of insurgency in Balochistan from 1948 to 2018 [32]. UCDP's version 26.1 replication release covers data through 2025, while its armed-conflict file includes Balochistan conflict-year records across multiple historical periods [28] [29]. Its actor dataset provides stable identifiers for Baloch conflict organizations [30]. Together, these sources support a historical sequence with changing organizations and periods; they do not support flattening seven decades into a single present-tense entity.

The public history should be organized around eras, claims, actors, state responses, and evidence coverage. Each era can display the date range supported by a source, the organizations represented in a versioned actor registry, whether the period meets a dataset's armed-conflict definition, and the limits of the source. Academic narrative can explain historical evolution [32]. UCDP can provide versioned identifiers and conflict-year structure [28] [29] [30]. Human Rights Watch can provide attributed rights context, but it should not be converted into an event total or treated as an adjudication of every disputed allegation [31].

Rights context is essential rather than peripheral. Human Rights Watch's 2026 Pakistan chapter documents continuing rights concerns and attributes disputed figures to their original monitors [31]. That attribution practice is a useful product requirement. When figures differ, Palimpsest should store the originating organization, definition, reference period, and confidence or dispute note. NarcoScope should render the disagreement, not average incompatible counts. A blank or withheld value should say why it is unavailable.

Infrastructure and conflict can be linked at a policy level without becoming tactical intelligence. USIP's 2017 report treated the insurgency as a major obstacle to Gwadar development [11]. The official CPEC progress page identifies a large energy project in Hub, Balochistan [2], while official Gwadar sources list ports, free-zone enterprises, and related development projects [17] [18] [19]. These observations justify examining whether project geography, institutional legitimacy, local service delivery, and insecurity appear together in public analysis. They do not justify identifying vulnerable facilities, predicting routes, ranking individuals, or showing rapidly updating incident coordinates.

The safe publication unit should consequently be aggregate and delayed. Suitable outputs include conflict-year presence, era-level actor continuity, province- or country-level annual totals where licensing allows, uncertainty bounds, methodology notes, and links to exact dataset versions. Unsuitable outputs include person dossiers, unit movements, tactical access routes, facility vulnerabilities, or a map intended to guide action in the field. Coordinates present in a source should be quarantined before any transformation reaches the public API.

Licensing reinforces this design. UCDP states that all its datasets are free and licensed under CC BY 4.0 [7]. Its API is RESTful and returns JSON [8], but a versioned bulk artifact may still be preferable for reproducibility because it can be hashed, stored with a codebook, and reprocessed deterministically. ACLED's Balochistan enhancement notes that data from January 2022 onward have been incorporated [9], yet its EULA controls use and redistribution [10]. The correct initial posture is therefore publishable UCDP aggregates plus ACLED source metadata or links until the exact ACLED use case is confirmed as licensed.

This boundary does not make the product shallow. It makes the historical analysis more reliable. Readers can follow how classifications, actors, source coverage, and rights reporting change over time without being exposed to a misleading “live intelligence” aesthetic. The system can deeply cover political economy, institutional grievance, project narratives, aggregate violence, and data disagreement while refusing the details most likely to create harm or be misinterpreted as operational truth.

Finding 5: Palimpsest should preserve evidence; NarcoScope should publish a controlled projection

The integration should be asymmetric. Palimpsest is the evidence ledger: it records source identity, retrieved bytes or a content hash, observed and published dates, evidence spans, project aliases, revisions, rights, and analytical claims. NarcoScope is the public query and narrative layer: it consumes a versioned release from Palimpsest, applies safety and licensing transforms, and exposes only approved aggregates. Copying the same raw records into both products would create two authorities, two freshness clocks, and an avoidable risk of leaking material that was valid for internal analysis but not public redistribution.

The source registry is the first shared primitive. A source has a stable identifier, canonical locator, title, author or issuing organization, edition or dataset version, source type, metadata-verification state, retrieval time, content hash, and rights state. The present research already demonstrates the value of stable identity: the source display numbers [1] through [34] are convenient for readers, while machine relationships should use immutable source identifiers. If bibliography order changes, a claim remains attached to the same source.

An evidence record is smaller than a source. It contains a quote, data point, or tightly bounded paraphrase; a page, table, row, endpoint, or section locator; capture time; and extraction method. A claim then links one or more evidence records and declares whether it is factual, synthetic,

recommended, or speculative. This prevents a long report or dataset landing page from being cited as if it supported every nearby sentence. It also permits a later reviewer to replace one weak evidence span without rewriting project identity.

Rights require a state machine. `publishable` means the intended transform is allowed with required attribution. `link_only` permits metadata and a link but no copied data. `restricted` means access may be lawful for a user but redistribution is not authorized. `unknown` means the evidence is quarantined until terms are reviewed. `unavailable` covers inaccessible sources without turning them into negative observations. UCDP's CC BY 4.0 declaration and the World Bank trade-cost dataset's CC BY 4.0 classification support publishable derived aggregates with attribution [7] [14]. The ADB outlook's CC BY 3.0 IGO status similarly supports reuse subject to its terms [27]. ACLED's EULA warning requires a more cautious per-use determination [10].

Freshness also needs multiple clocks. `source_published_at` says when an issuer released material. `observation_start` and `observation_end` define what period a number describes. `retrieved_at` records when Palimpsest observed the source. `transformed_at` records when the publication artifact was generated. `deployed_at` records when NarcoScope began serving it. The World Bank's biannual Myanmar monitor series [5] and UCDP's 26.1 release covering through 2025 [28] make clear why a newly retrieved page can still contain an older observation period. A single “updated today” badge would be false precision.

Project identity should be versioned independently of project observations. Names change, translations differ, and one program can contain multiple assets. A canonical project object therefore holds aliases, geography at a safe administrative level, sector, program memberships, and related agreements. Status, finance, throughput, and outcomes live in dated observations linked to sources. The official Gwadar portfolio, airport, free-zone enterprise count, concession, and CPEC priority plan can then relate to a Gwadar project family without being merged into one record [16] [17] [18] [19] [22].

The Palimpsest-to-NarcoScope release should be a signed or hashed manifest rather than an implicit database read. It would state schema version, source-set hash, record counts by publication state, maximum observation date, build time, transform version, and excluded-record counts by reason. NarcoScope would reject a release with a schema mismatch, invalid hash, missing rights decision, or prohibited fields. A successful HTTP response would prove only availability; the manifest and sampled evidence links would prove what release is actually being served.

Public APIs should make uncertainty queryable. A corridor endpoint could return projects grouped by country, program, sector, and lifecycle state; an economics endpoint could return indicators with units, observation periods, and source lineage; a conflict-context endpoint could return year-level aggregates and methodology; and an evidence endpoint could show only publishable citations. Every endpoint should return `data_status`, `rights_status`, `freshness_status`, and `limitations`. If no safe value exists, the response should be a typed unavailable state rather than `0`, an empty array with no explanation, or a fabricated estimate.

The user interface should follow the same model. Each headline number needs a visible source, reference period, publication date, and methodology. Official claims should be labeled “officially reported,” model results “modeled,” and observed series “measured” only when the underlying source supports that description. Contradictions should appear side by side. A project with strong documentation but no outcome evidence should look incomplete, not unsuccessful. A source with uncertain reuse rights should remain discoverable as metadata without silently appearing in charts.

Finally, integration quality should be tested against failure cases, not only happy paths. Tests should ensure that restricted rows never reach public artifacts; a stale observation cannot inherit a fresh retrieval date; a project alias cannot create duplicate totals; a missing value cannot become zero; a source revision changes the release hash; and conflict coordinates or person-level fields are rejected. These controls are what make a large system trustworthy enough for others to rely on deeply.

SYNTHESIS & INSIGHTS

Five patterns emerge from the evidence. First, every major topic is temporal. CPEC plans date from 2017 [22], Gwadar's concession transfer dates to 2013 [19], Pakistan's macro figures refer to FY25 [20], Myanmar's institutional interruption begins in February 2021 [26], and UCDP's current release covers through 2025 [28]. A product organized only around entities will mislead; it must organize entity observations along explicit timelines.

Second, official detail and independent evaluation are complements. Official portals are often the best source for project names, administrative status, agreements, and portfolios [1] [2] [16] [17] [18] [19] [23] [24] [25]. Independent and multilateral sources are better suited to corridor models, finance-system context, macro monitoring, conflict definitions, and rights documentation [3] [4] [5] [7] [13] [20] [21] [28] [31] [32] [34]. Removing official sources would erase primary claims; treating them as neutral outcome evaluation would erase the distinction between assertion and verification.

Third, the hardest gaps are not solved by adding more rows. Economic additionality needs a causal design. Local benefit needs distributional measures. Project completion needs operating evidence. Conflict counts need consistent definitions and licensed publication. These are schema and methodology problems before they are volume problems. The most valuable next source is often the one that completes an evidence chain—such as a financing agreement or throughput series—rather than another article repeating an announcement.

Fourth, political and economic risk cannot be reduced to a single score without hiding judgment. The Pakistan macro evidence identifies external financing constraints [20] [21]; the historical literature and datasets establish a long, changing Balochistan conflict context [11] [28] [29] [30] [32]; human-rights reporting adds attributed concerns [31]. None supplies a common natural unit. A composite score would require weights that are political choices. The better product lets users examine dimensions separately and publishes any derived index with its formula and sensitivity.

Fifth, the architecture itself can improve research behavior. A claim ledger encourages writers to make atomic statements. A rights state prevents an attractive but unlicensed dataset from leaking into publication. Separate freshness clocks expose stale observations. A release manifest lets users verify that the public service matches the reviewed evidence. These are not backend niceties; they are how product design reduces epistemic error.

The principal novel insight is that Palimpsest and NarcoScope should optimize different forms of completeness. Palimpsest seeks evidentiary completeness: every known source, version, contradiction, and gap that can lawfully be retained. NarcoScope seeks decision completeness: enough safe, current, intelligible context for a reader to understand a question and its uncertainty. The transformation between them is therefore a governed publication boundary, not a synchronization job.

A second insight is that “ground-level change” is best supported through accountable questions rather than sweeping verdicts. Users can ask what an authority promised, what was financed, what was built, what appears to operate, what changed in trade or service measures, who bears risk, and what evidence is missing. Those questions can drive reporting, oversight, academic replication, and policy debate. A dashboard that answers them with provenance is more consequential than a larger dashboard that collapses them into a confident narrative.

LIMITATIONS & CAVEATS

The evidence base is broad enough to define architecture and initial coverage, but not to establish an exhaustive project inventory or causal evaluation. It contains 34 registered sources and one bounded evidence span per source. That is a strong provenance skeleton, not a finished country database. Many official pages contain richer tables than the persisted span represents; many contracts, disbursement schedules, procurement records, tariff determinations, environmental assessments, cargo statistics, and local socioeconomic indicators have not yet been captured. The report therefore avoids claiming total CPEC or corridor investment, total Chinese debt, realized port throughput, or a measured welfare return.

The sources also have asymmetric authority. Pakistan and Myanmar government pages can establish what an authority publishes, including project priorities, counts, agreements, and declared status [1] [2] [16] [17] [18] [19] [22] [23] [24] [25]. They cannot independently validate their own impact. AidData offers global project coverage and a transparent methodology [4] [13] [34], but global coverage does not guarantee complete or current records for every Pakistan or Myanmar project. The World Bank corridor estimate is modeled [3]. Macro sources describe national conditions [5] [15] [20] [21] [26] [27], not project-specific causation.

Conflict evidence has definitional and ethical limits. UCDP's armed-conflict and actor datasets support versioned historical structure [28] [29] [30], but a dataset definition does not capture every form of political contention or harm. Human Rights Watch documents rights concerns through a different methodology [31]. ACLED's coverage is broader in some dimensions, yet its own material discusses

enhancements and limitations, and its EULA governs reuse [9] [10] [12]. Counts across these systems should not be compared until units, inclusion criteria, temporal coverage, and publication rights are aligned.

Counterevidence Register

The principal apparent contradiction in Myanmar is between continuing official project institutions and interrupted development activity. Official sources describe Kyaukphyu cooperation and agreements [23] [24], while ADB says its public-sector disbursements and new contracts have been on hold since February 2021 [26]. The statements concern different institutions and dimensions, so neither cancels the other. The impact is significant for status design but does not justify a conclusion that Kyaukphyu is either fully progressing or cancelled.

The principal apparent contradiction in Pakistan is between official progress evidence and macroeconomic vulnerability. Authorities list completed and ongoing Gwadar assets [16] [17], while the State Bank and IMF describe large external liabilities and continued dependence on policy execution and external financing [20] [21]. These statements can all be true. Project completion does not resolve macro risk, and macro risk does not prove that an individual asset failed. The impact is moderate for project status and significant for any claim about national economic benefit.

The corridor model creates a third tension. The World Bank estimated a potential travel-time reduction from completed corridors [3], yet the present ledger lacks realized operating and outcome evidence for Gwadar and Kyaukphyu. The model is not counterevidence to underperformance, nor is missing outcome evidence counterevidence to the model. The correct conclusion is that expected mechanism and observed result are not yet connected in this evidence set.

Known gaps and uncertainty

Priority gaps include project-level financing terms; disbursement and repayment schedules; procurement and beneficial-ownership records where lawfully public; asset commissioning and utilization; port cargo and airport traffic; power tariffs and reliability; land, environment, resettlement, and local-service outcomes; and subnational economic measures. Myanmar additionally needs reliable status evidence that can survive severe institutional discontinuity. Balochistan needs broader historical scholarship, official legal records, civil-society documentation, and comparable aggregate series without exposing people or operational detail.

Rights can change or be use-specific. A source labeled public does not automatically permit bulk redistribution, and a Creative Commons license still carries attribution and sometimes other obligations. The system must store the exact terms or a reviewed rights decision for the intended transform. This report's recommendation to hold ACLED data to metadata or link-only publication is precautionary, based on the EULA notice rather than a legal opinion [10].

RECOMMENDATIONS

Immediate implementation

First, establish the shared evidence contract before adding more collectors. The contract should define stable source, evidence, claim, project, organization, agreement, finance observation, status observation, indicator observation, geography, and release-manifest objects. Every observation needs a source identifier, evidence locator, observation period, retrieval time, rights state, freshness state, and confidence or dispute state. Project aliases and program relationships must be many-to-many so that Gwadar, CPEC, and individual assets do not become duplicate totals.

Second, implement rights and safety as fail-closed gates. UCDP and the World Bank trade-cost data can begin in publishable workflows under their stated CC BY 4.0 terms [7] [14]. ADB material can be processed according to its CC BY 3.0 IGO terms [27]. ACLED should remain metadata or link-only until the precise collection, storage, transformation, and redistribution use is reviewed against its EULA [10]. Conflict transformations must remove coordinates, free-text incident narratives, person fields, and tactical detail before a NarcoScope artifact exists.

Third, build the first country modules around evidence already strong enough to test the model. The Pakistan module should ingest the official CPEC and Gwadar inventories, planning records, Pakistan Economic Survey structure, State Bank macro series, and IMF context [1] [2] [15] [16] [17] [18] [19] [20] [21] [22]. The Myanmar module should ingest official Kyaukphyu and special-economic-zone assertions, World Bank monitor metadata, and ADB context [5] [23] [24] [25] [26] [27]. Every collector must store raw-content hashes or revision identifiers so changed source pages create new observations rather than overwriting history.

Fourth, add a versioned aggregate conflict pipeline. Start with UCDP 26.1 bulk releases and codebooks, using the licensed download and replication artifacts [7] [28] [29] [30]. Produce only year- and safe administrative-level aggregates, actor continuity summaries, uncertainty where supplied, and methodology links. Publish the dataset version and release hash on every response. Do not combine UCDP and ACLED counts until a documented crosswalk explains definitions and rights.

Product integration

Fifth, define the Palimpsest release artifact and NarcoScope acceptance contract. The artifact should contain schema version, source-set hash, transform version, records by geography and rights state, maximum observation date, exclusion counts, and prohibited-field audit results. NarcoScope should verify the manifest and expose the exact release identifier on health, API, MCP, and user-interface surfaces. Rollback should select an earlier valid manifest, not rebuild mutable source data.

Sixth, ship four user journeys. The project dossier answers what was announced, agreed, financed, built, operated, and measured. The corridor economics view separates modeled mechanisms from observed indicators. The evidence explorer shows source, span, version, rights, and contradiction. The aggregate conflict-context view shows eras, definitions, and rights reporting without operational detail. Each journey should include an explicit “what is unknown” panel.

Seventh, measure reliability before reach. Useful release metrics include the share of public claims with direct evidence spans; sources with reviewed rights; projects with non-duplicated canonical identity; indicators with all freshness clocks; contradictions surfaced rather than overwritten; restricted records caught by the publication gate; and public responses whose manifest hash matches the reviewed release. Traffic and citation by others matter only after these trust measures remain healthy.

Research expansion

Eighth, complete economic chains rather than collecting indiscriminately. For each flagship asset, seek the governing agreement, financing commitment, disbursement, implementation evidence, operating measure, and outcome indicator. Add UN Comtrade through a bounded, cached collector that records exact query parameters and respects endpoint limits [33]. Compare before-and-after indicators descriptively at first. Any causal claim should require a predeclared method, counterfactual, robustness checks, and a claim-specific evidence review.

Ninth, establish a recurring challenge process. Domain reviewers should be able to flag project duplicates, stale status, disputed figures, unsafe granularity, and rights errors. Corrections should create new evidence and release versions while preserving the old record. A public changelog should distinguish source revision, parser correction, analytical reinterpretation, and rights withdrawal.

The primary recommendation is to treat breadth as the result of a stable, testable evidence pipeline. Adding Pakistan, Myanmar, Gwadar, Kyaukphyu, BRI economics, and Balochistan in one unversioned dataset would create apparent scale. Building them as rights-gated observations connected through shared identity and released through a verifiable publication boundary creates a system that other people can responsibly use.

METHODOLOGY APPENDIX

Research process

This ultradeep research run began by defining a bounded question: how to cover Chinese economic activity, BRI infrastructure and economics, and aggregate conflict context in Pakistan and Myanmar without claiming exhaustive observation or publishing operational intelligence. Retrieval was decomposed into global BRI economics, development-finance datasets, Pakistan planning and macroeconomics, Gwadar institutions, Myanmar and Kyaukphyu institutions, conflict-data methods and rights, human-rights context, and future trade measurement. Web search was used because the preferred multi-provider search command was unavailable during the run.

Sources were registered before synthesis. Each of the 34 source records has a stable identifier derived from a canonical locator. Registration order produced display numbers [1] through [34], which are used only for reading. The evidence store contains 34 bounded spans—one quote, data point, or paraphrase per registered source—with a source identifier, locator, retrieval query, evidence type, and capture time. Claims in the separate ledger link to stable source and evidence identifiers rather than to display order.

Triangulation was performed by claim family. Official project statements were compared with planning, macroeconomic, or independent methodological sources where available. Modeled corridor economics was kept separate from observed project performance. Project status was separated from finance and macro context. Conflict history was checked across academic narrative, UCDP versioned data, ACLED methodology, and rights reporting. Contradictions were preserved whenever the observations concerned different dimensions or time periods.

The outline changed after retrieval. Rather than organizing the report only by country, the evidence supported a leading section on conditional BRI economics and a dedicated architecture finding. This adjustment was necessary because the same identity, rights, freshness, and status problems recur across Pakistan, Myanmar, and conflict data. No new section was added without at least one persisted supporting source.

Source profile and credibility

The registry contains government and public-authority pages, multilateral and central-bank material, academic datasets and methodology, documentation, a policy report, and human-rights reporting. It spans foundational material from 2017 through sources registered as current in 2026. Temporal diversity is valuable but creates a risk: older analysis is retained as historical evidence, never silently presented as current. Official-source credibility is considered high for the fact that an authority made a statement and lower for independent validation of its consequences. Dataset credibility depends on documented definitions, versioning, licensing, and reproducibility.

No single numeric credibility score was assigned because doing so would hide claim-dependent authority. The Gwadar Port Authority is a primary source for its published project list [17]; it is not automatically the strongest source for realized welfare impact. UCDP is a reproducible source for observations under its conflict definitions [7] [28] [29] [30]; it is not a complete measure of every political grievance. This role-based assessment is more informative than one source-wide rank.

Verification approach

Every factual statement that materially affects the findings is cited immediately. The bibliography is generated from the source registry and retains the registered title, year, author where available, and canonical URL. Claims are atomic: a model estimate, official status, finance observation, rights statement, or dataset-coverage fact is not bundled with a broader conclusion. Automated validators

check required sections, citation continuity, bibliography completeness, placeholder text, HTML coverage, and deterministic claim-to-evidence lexical and numeric support. Automated support is a floor, not a substitute for manual semantic review.

The rights review is use-specific. A public or Creative Commons classification is recorded as evidence of a publication path, subject to the actual license conditions. An EULA notice triggers quarantine or link-only treatment until the intended transformation is reviewed. Safety review applies independently of copyright: even licensed conflict data can be aggregated or withheld if detailed publication could expose people or enable operational use.

Claims-Evidence Table

Claim ID	Major claim	Evidence class	Supporting sources	Confidence
C1	BRI corridor travel-time improvement is a modeled conditional effect, not a realized Gwadar or Kyaukphyu return.	Multilateral model	[3] [14]	High for classification; low for local realized outcome
C2	Chinese development-finance sources require dataset- and version-specific counts.	Dataset and methodology	[4] [13] [34]	High
C3	Official Gwadar materials describe multiple distinct assets, agreements, and portfolio measures.	Government and port-authority records	[1] [2] [16] [17] [18] [19] [22]	High for attributed claims
C4	Pakistan project evidence must be contextualized without attributing the entire national external position to CPEC.	Central bank and IMF context	[20] [21]	High for macro context; attribution not established
C5	Myanmar project institutions and development interruption can coexist as different status dimensions.	Government and multilateral records	[23] [24] [25] [26]	Medium-high
C6	Balochistan requires a multi-period, multi-actor historical model.	Academic history and versioned conflict data	[28] [29] [30] [32]	Medium-high
C7	Conflict publication must be rights-gated and method-aware.	Dataset licenses, methods, and EULA	[7] [9] [10] [12]	High
C8	Palimpsest should preserve full lineage while NarcoScope publishes a coarser approved projection.	Cross-source synthesis	[5] [7] [10] [14] [27] [28]	Medium; architectural recommendation

Reproducibility and update policy

The research folder preserves `run_manifest.json`, `sources.jsonl`, `evidence.jsonl`, `claims.jsonl`, the Markdown report, and generated presentation artifacts. A future update should append new source and evidence identities, preserve prior releases, and record why a claim changed. Re-running a collector without a source revision should not manufacture a new observation. Reordering the report may change display citations, but stable source relationships remain intact.

The final quality boundary is explicit: this report is a detailed design and initial synthesis, not a legal opinion, a complete infrastructure census, a causal economic evaluation, or a live security product. Its confidence is high on the proposed evidence architecture and on the narrowly quoted registered facts, medium on cross-source policy synthesis, and deliberately unassigned where required source chains are missing.

BIBLIOGRAPHY

- [1] Government of Pakistan, CPEC Secretariat (2026). "Gwadar Projects Under CPEC". CPEC Authority. <https://cpec.gov.pk/gwadar> (Retrieved: 2026-08-26)
- [2] Government of Pakistan, CPEC Secretariat (2026). "CPEC Projects Progress Update". CPEC Authority. <https://cpec.gov.pk/progress-update> (Retrieved: 2026-08-26)
- [3] World Bank (2019). "Belt and Road Economics: Opportunities and Risks of Transport Corridors". World Bank. <https://documents1.worldbank.org/curated/en/715511560787699851/pdf/Main-Report.pdf> (Retrieved: 2026-08-26)
- [4] AidData (2023). "Global Chinese Development Finance Dataset, Version 3.0". AidData. <https://www.aiddata.org/data/aiddatas-global-chinese-development-finance-dataset-version-3-0> (Retrieved: 2026-08-26)
- [5] World Bank (2026). "Myanmar Economic Monitor Reports". World Bank. <https://www.worldbank.org/en/country/myanmar/publication/myanmar-economic-monitor-reports> (Retrieved: 2026-08-26)
- [6] International Monetary Fund (2026). "Pakistan and the IMF". IMF. <https://www.imf.org/en/countries/pak> (Retrieved: 2026-08-26)
- [7] Uppsala Conflict Data Program (2026). "UCDP Dataset Download Center". Uppsala University. <https://ucdp.uu.se/downloads/> (Retrieved: 2026-08-26)
- [8] Uppsala Conflict Data Program (2026). "UCDP Application Programming Interface". Uppsala University. <https://ucdp.uu.se/apidocs/index.html> (Retrieved: 2026-08-26)
- [9] ACLED (2026). "ACLED Data Enhancements and Known Limitations". ACLED. <https://acleddata.com/methodology/overview-acleddata-data-enhancement-projects-and-known-data-limitations> (Retrieved: 2026-08-26)
- [10] ACLED (2026). "myACLED FAQs and EULA". ACLED. <https://acleddata.com/myacled-faqs> (Retrieved: 2026-08-26)
- [11] United States Institute of Peace (2017). "The China-Pakistan Economic Corridor". USIP. <https://www.usip.org/sites/default/files/2017-10/pw135-the-china-pakistan-economic-corridor.pdf> (Retrieved: 2026-08-26)
- [12] ACLED (2026). "ACLED Methodology on Myanmar". ACLED. <https://acleddata.com/methodology/acleddata-methodology-myanmar> (Retrieved: 2026-08-26)
- [13] Samantha Custer et al. (2023). "Tracking Chinese Development Finance: TUFF 3.0 Methodology". AidData. https://docs.aiddata.org/ad4/pdfs/AidData_TUFF_methodology_3_0.pdf (Retrieved: 2026-08-26)

- [14] World Bank (2019). "Belt and Road Initiative Trade Costs Database". World Bank Data Catalog. <https://datacatalog.worldbank.org/search/dataset/0038706/belt-and-road-initiative-trade-costs-database> (Retrieved: 2026-08-26)
- [15] Government of Pakistan Ministry of Finance (2025). "Pakistan Economic Survey 2024-25". Ministry of Finance. https://finance.gov.pk/survey_2025.html (Retrieved: 2026-08-26)
- [16] Government of Pakistan Planning Commission (2025). "Pakistan Annual Plan 2025-26". Planning Commission. https://www.pc.gov.pk/uploads/annualplan/Annual_Plan_2025_26.pdf (Retrieved: 2026-08-26)
- [17] Gwadar Port Authority (2026). "Gwadar Port Authority Development Projects". Gwadar Port Authority. <https://www.gwadarport.gov.pk/project.aspx> (Retrieved: 2026-08-26)
- [18] Gwadar Port Authority (2026). "Gwadar Port Free Zone Development Status". Gwadar Port Authority. <https://www.gwadarport.gov.pk/freezone.aspx> (Retrieved: 2026-08-26)
- [19] Gwadar Port Authority (2026). "Gwadar Port Master Plan". Gwadar Port Authority. <https://gwadarport.gov.pk/masterplan.aspx> (Retrieved: 2026-08-26)
- [20] State Bank of Pakistan (2025). "The State of Pakistans Economy Annual Report 2024-25". State Bank of Pakistan. <https://www.sbp.org.pk/reports/annual/aarFY25/Complete.pdf> (Retrieved: 2026-08-26)
- [21] International Monetary Fund (2025). "Pakistan Second Review Under the Extended Fund Facility". IMF. <https://www.imf.org/en/publications/cr/issues/2025/12/11/pakistan-second-review-under-the-extended-arrangement-under-the-extended-fund-facility-572472> (Retrieved: 2026-08-26)
- [22] Government of Pakistan and Government of the People's Republic of China (2017). "Long Term Plan for China-Pakistan Economic Corridor". Planning Commission of Pakistan. <https://pc.gov.pk/uploads/cpec/LTP.pdf> (Retrieved: 2026-08-26)
- [23] Myanma Port Authority (2026). "Kyaukphyu Deep Sea Port". Myanma Port Authority. <https://www.mpa.gov.mm/ports/kyaukphyu-deep-sea-port/> (Retrieved: 2026-08-26)
- [24] Myanmar Ministry of Information (2025). "Myanmar and CITIC Coordinate Kyaukpyu SEZ Implementation Conditions". Ministry of Information. <https://www.moi.gov.mm/moi%3Aeng/index.php/news/16809> (Retrieved: 2026-08-26)
- [25] Directorate of Investment and Company Administration (2026). "Why Invest in Myanmar". DICA. <https://www.dica.gov.mm/why-invest-in-myanmar/> (Retrieved: 2026-08-26)
- [26] Asian Development Bank (2026). "Myanmar Economy and Forecasts". ADB. <https://www.adb.org/where-we-work/myanmar/economy> (Retrieved: 2026-08-26)
- [27] Asian Development Bank (2025). "Asian Development Outlook April 2025". ADB. <https://www.adb.org/publications/asian-development-outlook-april-2025> (Retrieved: 2026-08-26)
- [28] Uppsala Conflict Data Program (2026). "UCDP Replication Datasets". Uppsala University. https://ucdp.uu.se/downloads/replication_data.html (Retrieved: 2026-08-26)
- [29] Uppsala Conflict Data Program (2026). "UCDP PRIO Armed Conflict Dataset 26.1 CSV". Uppsala University. <https://ucdp.uu.se/downloads/ucdprio/ucdp-prio-acd-261-csv.zip> (Retrieved: 2026-08-26)
- [30] Uppsala Conflict Data Program (2026). "UCDP Actor Dataset 26.1 CSV". Uppsala University. <https://ucdp.uu.se/downloads/actor/ucdp-actor-261-csv.zip> (Retrieved: 2026-08-26)
- [31] Human Rights Watch (2026). "World Report 2026 Pakistan". Human Rights Watch. <https://www.hrw.org/world-report/2026/country-chapters/pakistan> (Retrieved: 2026-08-26)
- [32] Noor Saeed Khattak, Alia Naseer, and Azeem Gul (2022). "History and Evolution of Insurgency in Balochistan 1948 to 2018". Global Political Review. [https://doi.org/10.31703/gpr.2022\(vii-iii\).04](https://doi.org/10.31703/gpr.2022(vii-iii).04) (Retrieved: 2026-08-26)
- [33] United Nations Statistics Division (2026). "UN Comtrade API Documentation". United Nations. <https://uncomtrade.org/docs/un-comtrade-api/> (Retrieved: 2026-08-26)
- [34] Seth Goodman et al. (2024). "AidData Geospatial Global Chinese Development Finance Dataset". AidData. <https://www.aiddata.org/publications/aiddatas-geospatial-global-chinese-development-finance-dataset> (Retrieved: 2026-08-26)